

1. Explain the problem or question that the group explored in a paragraph or two;

Before starting to work on this project, there are two things to do: read the E1000 manual and the driver code. From the manual, we can understand roughly the driver's data structure, like descriptor, command, status...etc., and the driver's transmitting and receiving flow.

Afterward, we have the basic knowledge to the code. We start from the driver code. There are two components to be finished, `e1000_transmit()` and `e1000_rcv()`. We can complete them via the explanation on MIT website.

2. Explain the key aspects of the work you did;

To make the driver act, I need to finish two components left by MIT team: `e1000_transmit()` and `e1000_rcv()`. They are used to set hardware before transmitting and after receiving message to/from server.

```

95  int
96  e1000_transmit(struct mbuf *m)
97  {
98      //
99      // Your code here.
100     //
101     // the mbuf contains an ethernet frame; program it into
102     // the TX descriptor ring so that the e1000 sends it. Stash
103     // a pointer so that it can be freed after sending.
104     //
105     //printf("Suwei: something need to send here.\n");
106     acquire(&e1000_lock);
107     uint32_t transit_tail = regs[E1000_TDT]; //get available index
108     if((tx_ring[transit_tail].status & E1000_TXD_STAT_D0) == 0) //check if available tx desc is ready.
109     {
110         release(&e1000_lock);
111         return -1; //return error, because not ready
112     }
113     if(tx_mbufs[transit_tail] != 0)
114     {
115         mbuf_free(tx_mbufs[transit_tail]); //free mbuf from available tx descriptor
116         tx_ring[transit_tail].addr = (uint64_t)m->head;
117         tx_ring[transit_tail].length = m->len;
118         tx_ring[transit_tail].cmd = E1000_TXD_CMD_EOP | E1000_TXD_CMD_RS; //set cmd to hw
119         tx_mbufs[transit_tail] = m;
120         regs[E1000_TDT] = (transit_tail + 1) % TX_RING_SIZE; //add tail
121         release(&e1000_lock);
122         return 0;
123     }

```

From `e1000_transmit()`, we will get a buffer, mbuf, containing a couple of things: message length, content, IP, and port. We need to put information into corresponding registers and then issue a command afterward. Hardware will push the data to the server we specified. Finally, we need to increase the tail by 1 to tell the hardware the next available space for the next transmission.

```

125 static void
126 e1000_recv(void)
127 {
128     //
129     // Your code here.
130     //
131     // Check for packets that have arrived from the e1000
132     // Create and deliver an mbuf for each packet (using net_rx()).
133     //
134     //printf("Suwei: something need to receive here.\n");
135     uint32 receive_tail = (regs[E1000_RDT] + 1) % RX_RING_SIZE; //get received tail
136     struct rx_desc *rx_desc = &rx_ring[receive_tail]; //get received descriptor
137     struct mbuf *m = rx_mbufs[receive_tail];
138     while(rx_desc->status & E1000_RXD_STAT_DD)
139     {
140         m->len = rx_desc->length;
141         net_rx(m);
142         m = mbufalloc(0);
143         rx_mbufs[receive_tail] = m; //update new sapce for mbuf
144         rx_desc->addr = (uint64)m->head;
145         rx_desc->status = 0; //clear status bits
146         receive_tail = (receive_tail + 1) % RX_RING_SIZE; //add tail
147         rx_desc = &rx_ring[receive_tail];
148         m = rx_mbufs[receive_tail];
149         //need loops in this function to check the tail.
150     }
151     regs[E1000_RDT] = (receive_tail - 1) % RX_RING_SIZE;
152 }
153

```

In `e1000_recv()`, we will need to handle the flow after receiving the message from server. Only two things that need to be finished: putting mbuf into `net_rx()` and clearing space for following received package. In `net_rx()`, driver will find a corresponding socket for received package, as shown below. Then, clear move can have a space for upcoming message.

```

161 void
162 sockrecvd(struct mbuf *m, uint32 raddr, uint16 lport, uint16 rport)
163 {
164     //
165     // Find the socket that handles this mbuf and deliver it, waking
166     // any sleeping reader. Free the mbuf if there are no sockets
167     // registered to handle it.
168     //
169     struct sock *si;
170
171     acquire(&lock);
172     si = sockets;
173     while (si) {
174         if (si->raddr == raddr && si->lport == lport && si->rport == rport)
175             goto found;
176         si = si->next;
177     }
178     release(&lock);
179     mbuffree(m);
180     return;
181
182 found:
183     acquire(&si->lock);
184     mbufq_push(&si->rbuf, m);
185     wakeup(&si->rbuf);
186     release(&si->lock);
187     release(&lock);
188 }

```

After finishing the two components, the driver can pass the test script, and we can start doing some improvement. I found the they used `copyin()/copyout()` as the communication between user and kernel, which might be the bottleneck for the driver. Hence I created another function, `mbufalloc_suwei`, to map mbuf into user memory space (shown below). Hoping to improve the performance by reducing memory movement during the kernel.

```

139 int
140 sockwrite(struct sock *si, uint64 addr, int n)
141 {
142     struct proc *pr = myproc(); //suwei comment
143     struct mbuf *m;
144
145     m = mbufalloc(MBUF_DEFAULT_HEADROOM); //suwei comment
146     //m = mbufalloc_suwei(MBUF_DEFAULT_HEADROOM, addr, n); //suwei add
147     //mbufput(m, n); //suwei add
148     if (!m)
149         return -1;
150     //if 1/suwei add
151     if (copyin(pr->pagetable, mbufput(m, n), addr, n) == -1) {
152         mbuffree(m);
153         return -1;
154     }
155     #endif
156     net_tx_udp(m, si->raddr, si->lport, si->rport);
157     return n;
158 }

```

```

64 //Allocates a packet buff, suwei add
65 //use user pa as driver space.
66 struct mbuf *
67 mbufalloc_suwei(unsigned int headroom, uint64 addr, int n)
68 {
69     struct proc *pr = myproc();
70     pagetable_t pagetable = pr->pagetable;
71     uint64 va0, pa0;
72     struct mbuf *m;
73
74     va0 = PGROUNDUP(addr);
75     pa0 = walkaddr(pagetable, va0); //va is continuous, but pa is continuous in a 4K
76     //memset((void *)pa0, 0, PGSIZE); //suwei add
77     if (headroom > MBUF_SIZE)
78         return 0;
79     //m = kalloc(); //suwei comment
80     m = (struct mbuf *)pa0; //suwei add
81     if (m == 0)
82         return 0;
83     m->next = 0;
84     m->head = (char *)m->buf + headroom;
85     m->len = 0;
86     //memset(m->buf, 0, sizeof(m->buf)); //suwei comment
87     return m;

```

3. Explain what was interesting/exciting/promising about what you did;

The most exciting part for me was when the driver sent package to server and when performance got different after I modified it. It still has some improved space like read flow. We can create buffer space for the message and map it directly to user space to further reduce the memory operation. However, it might take more time to implement and debug.

4. Explain what limitations you ran into in the work; what questions couldn't you answer; what prevented you from answering them;

From the result shown below, single process test improved after modification. I used a longer message to test afterward and it was 11ms(improved) v.s. 29ms(original). However, multi-process test elapsed time increased after the improvement. I didn't figure out what they were waiting for. Especially when I increased the message length, they increased further. I guess it is probably related to the job in user space.

5. Result

The image contains four terminal screenshots. The top-left screenshot shows the command 'make qemu' being executed, followed by the output of 'nettests' which includes 'single-process write time period: 11' and 'multi-process write time period: 0'. The top-right screenshot shows a series of 'message from xv6!' messages. The bottom-left screenshot shows the command 'make qemu' being executed, followed by the output of 'nettests' which includes 'single-process write time period: 8' and 'multi-process write time period: 3'. The bottom-right screenshot shows a series of 'cpy msg bla bla bla!' messages.

Single process/ uptime()	T1	T2	T3	T4	T5
Original	11	12	13	15	14
Improved	8	9	15	11	15